COMP SCI 397
Course Project Final Report
Tracking State Changes in Procedural Text

Danilo Neves Ribeiro
dnr2876
daniloribeiro2021@
u.northwestern.edu

William Hancock
wwh7118
WilliamHancock2022@
u.northwestern.edu

Abstract

In this project we worked on the recent AI2 question-answering benchmark called ProPara. This data set is comprised of procedural text paragraphs covering different topics (e.g. photosynthesis) and the objective is to keep track of how state of entities involve through time (e.g. water or light gets absorbed by the plant). Our approach follows the ProGlobal model from the original ProPara paper. We show how our results compare to the published results and also compoile an error analysis on the results of the original paper.

1 Introduction

Answering questions about paragraphs that describe processes is still a challenging task for machine reading comprehension systems. This genre of text is pervasive (e.g. manuals, recipes, road safety rules, scientific protocols, etc.) and understanding them often requires keeping track of how the world's state evolve over time. For instance, consider the paragraph describing photosynthesis in Figure 1. If the system is asked the question: "Where is sugar produced?", it is expected to answer "In the leaf". To answer the question, the system needs to infer the state changes of each entity in the paragraph and the causality between such change events (which are often implicit, making this a challenging task). The dataset is further detailed in the following section.

2 Data Set

To evaluate our system, we use the ProPara procedural text benchmark which contains 488 crowdsourced paragraphs and 3100 sentences total. This data set is comprised of procedural text paragraphs covering different topics together with a human annotated table that describes the state (location and existence) of entities in this paragraph. Figure 1 shows an instance of the training data which constitutes of a paragraph about photosynthesis and the annotated state change grid. The state change grid contains information about where an entity (e.g. water or light) is at each step. Note that "?" indicates the location is unknown, and "-" indicates the entity doesn't exist during that step.

		Participants:				
Paragraph (seq. of steps):		water	light	CO2	mixture	sugar
	state0	soil	sun	?	-	-
Roots absorb water from soil						
	state1	roots	sun	?	-	-
The water flows to the leaf.						
	state2	leaf	sun	?	-	-
Light from the sun and CO2 enter the leaf.						
	state3	leaf	leaf	leaf	-	-
The light, water, and CO2 combine into a mixture.						
	state4	-	-	-	leaf	-
Mixture forms sugar.						
	state5	-	-	-	-	leaf

Time

Figure 1: ProPara participant state change grid.

3 Model and Implementation

Attention in proglobal is more complicated than that of a simple seq2seq architecture, due to the fact that proglobal is more complex. A bidirectional lstm learns a hidden global state H for each paragraph. This hidden state is then passed to various submodules.

The architecture works as follows. For each sentence in each paragraph, learn information about a participating entity, e.g. water. First, a decision is made as to the state of the entity in a sentence, e.g. exists, notexists, or unknown. If exists, then infer the location of the entity, e.g. "water is in the ground".

The existential submodule arguably does not use attention, as it takes softmax(H) as input. On the other hand, the location submodule receives the full paragraph hidden state. We note that the visualizations show the value of every word in ev-

ery timestep. However, the model only considers words within each timestep (in other words it does not look ahead to future sentences to infer location).

4 Evaluation and Results

During testing the system will 4 categories of questions from the output state change grid:

1. What are the Inputs? That is, which participants existed before the procedure began
2. What are the Outputs? That is, which participants existed after the procedure ended?
3. What are the Conversions? That is, which participants were converted to which other participants?
4. What are the Moves? That is, which participants moved from one location to another?

Note that these questions are templated, meaning they can be deterministically answered using the output state grid. The evaluation code was made available at https://github.com/allenai/aristo-leaderboard/tree/master/propara by the AI2 team. The code outputs precision, recall and F1 score for each question category.

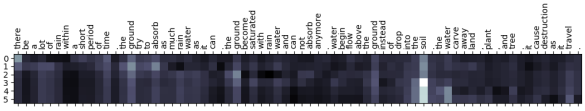

Figure 2: Rain

For "rain", the network infers the location "there" in the first sentence. This is an interesting artifact of this neural model It assumes that it has to tie a location to a word. "there" seems odd, but can be interpreted as technically correct. This is a limitation of the problem as it is set up, and not anything with the architecture itself.

The model infers that "rain" is in the ground in sentences 2 and 3. Sentence 4 is complicated, in that it is ambiguous. The water divides itself, now being both in and above the soil. The model actually infers that the water goes from the "ground" to the "soil", not realizing that this is the same thing. From a commonsense perspective however, the correct answer seems to be "above the ground". The negation of "instead" seems to be ignored by the model. It is clear that "soil" is overwhelmingly preferred by the attention mechanism. Perhaps because it is close in embedding space to "rain"?

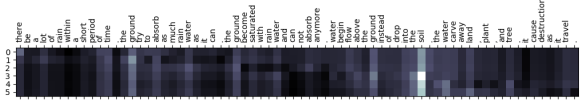

Figure 3: Plant

The model infers that "plant" appears in the last sentence, and is in the soil. I have no interpretation for this, as "soil" does not appear in the last sentence. Because this network is recurrent however, the overwhelming weight on "soil" must force the model to choose this location. "tree" (not shown) looks very identical to "plant", but the models infers that it is never introduced into a location.

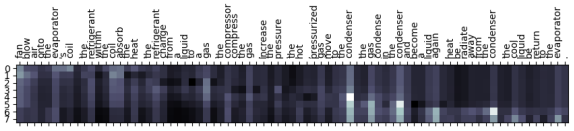

Figure 4: Heat

The model infers heat's path is "fan -¿ fan -¿ null -¿ condenser -¿ null -¿ condenser -¿ null" This thrashing is interesting, and it doesn't follow the narrative in the paragraph. The correct trajectory should be "air -¿ coil / gas / refrigerant". Clearly there is no commonsense here, the model seems to have a high prior for "condenser" and is overwhelmingly selecting this location. There is no commonsense of temporal stability either, in that it seems unreasonable for the location to fluctuate as it does.

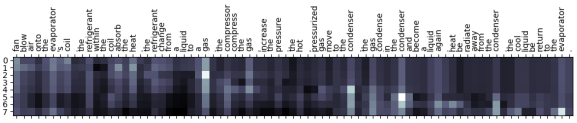

Figure 5: Liquid

To further illustrate this point, the model infers that liquid is at "gas" in sentence 3. This can almost certainly be attributed to a close embedding. There is no commonsense understanding that gas is not a container, nor that something cannot be liquid and gas.

5 Conclusion

It is clear that the problem of reasoning on a human level about these types of causal narratives is elusive. Can this model be said to be performing commonsense reasoning? Is this model useful in any sense? Can DL ever solve these problems? Is this a useful empirical evaluation?

The evaluation itself is oversimplified. This is almost certainly in order to make backprop more tractable while still appearing to be getting interesting results. The constraint on one span per transformation is evidence of this. This is simply unreasonable. An elaborate discussion on this matter can be found here: https://towardsdatascience.com/the-emperors-new-benchmarks-8fe8f170923b. Suffice it to say, the system will never grasp the complexity of the intended reasoning domain. This paper seems to show that neural networks can learn arbitrarily complex functions. That's it. The problem is that the function being learned in this situation isn't inherently useful (it will probably not generalize to other data of this sort).

How general is the ProGlobal model architecture? It does seem to be making some effort to be globally generally. Its assumptions are that there are objects, and their locations should be tracked. It is assuming some sort of persistence and continuation. In this sense it does seem to reasonably model a low level of causal understanding. Perhaps more deductive reasoners can be placed on top of this information to produce more intuitive and interesting results. It seems reasonable to question whether this will actually work, however. There seems to be a need for commonsense reasoning even at this low level in order to prevent some of the nonsensical conclusions.